

Coursera - Practical Machine Learning

Ali Magzari

8/30/2021

Introduction

This project consists of predicting the manner in which participants executed an exercise. The data is collected from accelerometers placed on the belt, forearm, arm, and dumbell of six participants.

For a complete explanation: <https://www.coursera.org/learn/practical-machine-learning/supplement/PvInj/course-project-instructions-read-first>

Methodology

This part covers the methods used to explore the data and choose the appropriate predictive algorithm.

Acquiring, cleaning and exploring data

Set specific seed for experiment replication purposes.

```
set.seed(500)

library(caret)

## Loading required package: lattice

## Loading required package: ggplot2

library(rpart)
library(rpart.plot)
library(rattle)

## Loading required package: tibble

## Loading required package: bitops

## Rattle: A free graphical interface for data science with R.
## Version 5.4.0 Copyright (c) 2006–2020 Togaware Pty Ltd.
## Type 'rattle()' to shake, rattle, and roll your data.
```

```
library(e1071)
```

Reading both training and testing data

```
data_train <- read.csv("C:/MY FILES/Rdirectory/pml-training.csv")
data_test <- read.csv("C:/MY FILES/Rdirectory/pml-testing.csv")
```

After looking at the data structure and reading carefully the variables, the following ought to be mentioned:

- The data frame contains 19622 observations of 160 variables.
- The desired variable to predict is named “classe” and takes 5 different values (A-E).
- The first 7 variables have no effect on “classe” and can therefore be omitted.

```
data_train <- data_train[, -c(1:7)]
```

```
data_test <- data_test[, -c(1:7)]
```

```
dim(data_train)
```

```
## [1] 19622 153
```

```
dim(data_test)
```

```
## [1] 20 153
```

Both sets have 153 variables, which is logical. The original sets enclosed 160 variables each. After removing 7 columns, we should be left with 153.

After careful inspection, it came to our attention that a lot of columns contain NA entries.

```
colSums(is.na(data_train))
```

```
##          roll_belt           pitch_belt           yaw_belt
##                  0                  0                  0
##      total_accel_belt kurtosis_roll_belt kurtosis_pitch_belt
##                  0                  0                  0
##      kurtosis_yaw_belt skewness_roll_belt skewness_roll_belt.1
##                  0                  0                  0
##      skewness_yaw_belt      max_roll_belt      max_pitch_belt
##                  0                 19216                 19216
##      max_yaw_belt        min_roll_belt        min_pitch_belt
##                  0                 19216                 19216
##      min_yaw_belt        amplitude_roll_belt amplitude_pitch_belt
##                  0                 19216                 19216
##      amplitude_yaw_belt    var_total_accel_belt      avg_roll_belt
##                  0                 19216                 19216
##      stddev_roll_belt        var_roll_belt        avg_pitch_belt
##                  19216                 19216                 19216
##      stddev_pitch_belt        var_pitch_belt        avg_yaw_belt
##                  19216                 19216                 19216
```

```

##      stddev_yaw_belt           var_yaw_belt           gyros_belt_x
##                      19216                  19216                  0
##      gyros_belt_y              gyros_belt_z           accel_belt_x
##                      0                      0                  0
##      accel_belt_y             accel_belt_z           magnet_belt_x
##                      0                      0                  0
##      magnet_belt_y            magnet_belt_z          roll_arm
##                      0                      0                  0
##      pitch_arm                 yaw_arm               total_accel_arm
##                      0                      0                  0
##      var_accel_arm            avg_roll_arm          stddev_roll_arm
##                      19216                  19216                  19216
##      var_roll_arm              avg_pitch_arm         stddev_pitch_arm
##                      19216                  19216                  19216
##      var_pitch_arm             avg_yaw_arm          stddev_yaw_arm
##                      19216                  19216                  19216
##      var_yaw_arm               gyros_arm_x          gyros_arm_y
##                      19216                  0                      0
##      gyros_arm_z               accel_arm_x          accel_arm_y
##                      0                      0                  0
##      accel_arm_z               magnet_arm_x         magnet_arm_y
##                      0                      0                  0
##      magnet_arm_z              kurtosis_roll_arm    kurtosis_pictch_arm
##                      0                      0                  0
##      kurtosis_yaw_arm          skewness_roll_arm   skewness_pitch_arm
##                      0                      0                  0
##      skewness_yaw_arm          max_roll_arm          max_picth_arm
##                      0                  19216                  19216
##      max_yaw_arm                min_roll_arm          min_pitch_arm
##                      19216                  19216                  19216
##      min_yaw_arm                amplitude_roll_arm  amplitude_pitch_arm
##                      19216                  19216                  19216
##      amplitude_yaw_arm          roll_dumbbell        pitch_dumbbell
##                      19216                  0                      0
##      yaw_dumbbell               kurtosis_roll_dumbbell kurtosis_picth_dumbbell
##                      0                      0                  0
##      kurtosis_yaw_dumbbell     skewness_roll_dumbbell skewness_pitch_dumbbell
##                      0                      0                  0
##      skewness_yaw_dumbbell     max_roll_dumbbell    max_picth_dumbbell
##                      0                  19216                  19216
##      max_yaw_dumbbell          min_roll_dumbbell   min_pitch_dumbbell
##                      0                  19216                  19216
##      min_yaw_dumbbell          amplitude_roll_dumbbell amplitude_pitch_dumbbell
##                      0                  19216                  19216
##      amplitude_yaw_dumbbell    total_accel_dumbbell var_accel_dumbbell
##                      0                      0                  19216
##      avg_roll_dumbbell         stddev_roll_dumbbell var_roll_dumbbell
##                      19216                  19216                  19216
##      avg_pitch_dumbbell        stddev_pitch_dumbbell var_pitch_dumbbell
##                      19216                  19216                  19216
##      avg_yaw_dumbbell          stddev_yaw_dumbbell  var_yaw_dumbbell
##                      19216                  19216                  19216
##      gyros_dumbbell_x          gyros_dumbbell_y     gyros_dumbbell_z
##                      0                      0                  0

```

```

##          accel_dumbbell_x      accel_dumbbell_y      accel_dumbbell_z
##                      0                  0                  0
##          magnet_dumbbell_x     magnet_dumbbell_y     magnet_dumbbell_z
##                      0                  0                  0
##          roll_forearm        pitch_forearm        yaw_forearm
##                      0                  0                  0
##          kurtosis_roll_forearm kurtosis_pitch_forearm kurtosis_yaw_forearm
##                      0                  0                  0
##          skewness_roll_forearm skewness_pitch_forearm skewness_yaw_forearm
##                      0                  0                  0
##          max_roll_forearm     max_pitch_forearm     max_yaw_forearm
##                      19216              19216              0
##          min_roll_forearm     min_pitch_forearm     min_yaw_forearm
##                      19216              19216              0
##          amplitude_roll_forearm amplitude_pitch_forearm amplitude_yaw_forearm
##                      19216              19216              0
##          total_accel_forearm   var_accel_forearm   avg_roll_forearm
##                      0                  19216              19216
##          stddev_roll_forearm   var_roll_forearm   avg_pitch_forearm
##                      19216              19216              19216
##          stddev_pitch_forearm  var_pitch_forearm   avg_yaw_forearm
##                      19216              19216              19216
##          stddev_yaw_forearm   var_yaw_forearm   gyros_forearm_x
##                      19216              19216              0
##          gyros_forearm_y       gyros_forearm_z    accel_forearm_x
##                      0                  0                  0
##          accel_forearm_y       accel_forearm_z    magnet_forearm_x
##                      0                  0                  0
##          magnet_forearm_y      magnet_forearm_z    classe
##                      0                  0                  0
##
```

Each corrupted column have 19216 NA entries (out of 19622)

The code above removes all NA columns and checks if all NA entries have been deleted.

```
data_train<- data_train[, colSums(is.na(data_train)) == 0]
```

```
any(colSums(is.na(data_train)) != 0)
```

```
## [1] FALSE
```

The same code block above should be applied to testing data.

```
data_test<- data_test[, colSums(is.na(data_test)) == 0]
```

```
any(colSums(is.na(data_test)) != 0)
```

```
## [1] FALSE
```

At this point, the datasets have the following dimensions:

```
dim(data_train)

## [1] 19622     86

dim(data_test)

## [1] 20 53
```

Preparing data

Data sets may contain predictors which value does not change dramatically across observations. These variables, usually called near-zero variance predictors are not informative and could therefore be deleted.

```
data_train <- data_train[, -nearZeroVar(data_train)]

dim(data_train)

## [1] 19622     53
```

The line above shows that 33 variables (86 - 53) were judged to be of near-zero variance, and were naturally deleted.

Time to partition our data(training) into training (70%) and validation (30%) sets.

```
data_part <- createDataPartition(data_train$classe, p = 0.7, list = FALSE)
data_train_p <- data_train[data_part, ]
data_valid_p <- data_train[-data_part, ]

nrow(data_train_p)

## [1] 13737

nrow(data_valid_p)

## [1] 5885
```

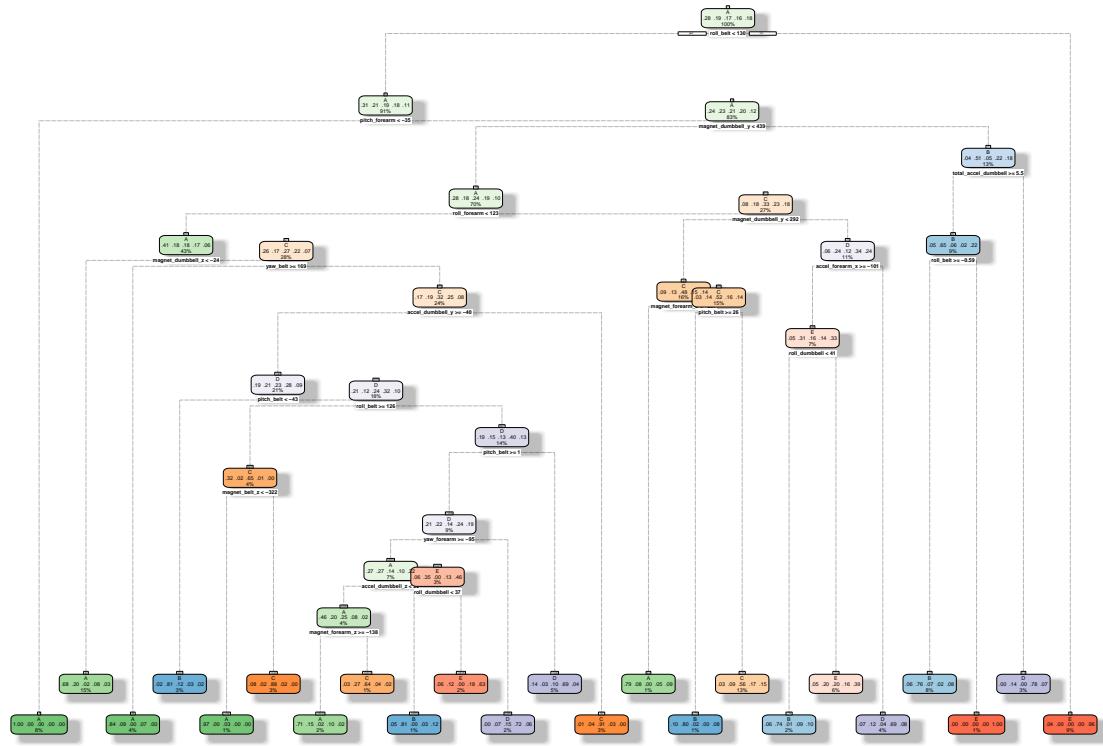
Model building

Two models will be constructed in this part: Decision trees and random forest.

Decision tree

```
model_tree <- rpart(classe ~ ., data = data_train_p, method = "class")
fancyRpartPlot(model_tree)

## Warning: labs do not fit even at cex 0.15, there may be some overplotting
```



Rattle 2021-Aug-31 23:25:20 21264

As shown above, the high number of leaf nodes makes it delicate to visualize the decision tree. Increasing the node size would only cause overlapping.

Let us now use the developed model to predict the variable “class” from the validation set.

```
model_tree_pred <- predict(model_tree, data_valid_p, type = "class")
confMat_tree <- confusionMatrix(model_tree_pred, factor(data_valid_p$class))
confMat_tree
```

```
## Confusion Matrix and Statistics
##
##             Reference
## Prediction   A    B    C    D    E
##           A 1498  227   27  106   37
##           B   37  630   44   21   63
##           C   36  114  835  148  122
##           D   49   92   61  615   57
##           E   54   76   59   74  803
##
## Overall Statistics
##
##                 Accuracy : 0.7444
##                 95% CI : (0.7331, 0.7555)
##     No Information Rate : 0.2845
##     P-Value [Acc > NIR] : < 2.2e-16
##
```

```

##                               Kappa : 0.6755
##
##  Mcnemar's Test P-Value : < 2.2e-16
##
##  Statistics by Class:
##
##                                Class: A Class: B Class: C Class: D Class: E
## Sensitivity                  0.8949   0.5531   0.8138   0.6380   0.7421
## Specificity                  0.9057   0.9652   0.9136   0.9474   0.9452
## Pos Pred Value                0.7905   0.7925   0.6653   0.7037   0.7533
## Neg Pred Value                0.9559   0.9000   0.9587   0.9304   0.9421
## Prevalence                    0.2845   0.1935   0.1743   0.1638   0.1839
## Detection Rate                 0.2545   0.1071   0.1419   0.1045   0.1364
## Detection Prevalence          0.3220   0.1351   0.2133   0.1485   0.1811
## Balanced Accuracy              0.9003   0.7592   0.8637   0.7927   0.8437

```

The confusion matrix displays that the decision tree model was able to predict the class in the validation set with an accuracy of 0.74%, leaving the out-of-sample error to be around 0.26%.

Random forest

This section will cover the construction of a random forest model.

```

control_rf <- trainControl(method = "cv", number = 3, verboseIter = FALSE)
model_rf <- train(classe ~ ., data = data_train_p, method = "rf", trControl = control_rf)

```

Let us now use the developed model to predict the variable “class” from the validation set.

```

model_rf_pred <- predict(model_rf, data_valid_p)
confMat_rf <- confusionMatrix(model_rf_pred, factor(data_valid_p$classe))
confMat_rf

```

```

## Confusion Matrix and Statistics
##
##                                Reference
## Prediction      A      B      C      D      E
##               A 1669     1      0      0      0
##               B     5 1136     2      1      1
##               C     0     2 1021    15      2
##               D     0     0     3 948     5
##               E     0     0     0     0 1074
##
## Overall Statistics
##
##                               Accuracy : 0.9937
##                               95% CI : (0.9913, 0.9956)
##  No Information Rate : 0.2845
##  P-Value [Acc > NIR] : < 2.2e-16
##
##                               Kappa : 0.992
##
##  Mcnemar's Test P-Value : NA

```

```

## 
## Statistics by Class:
## 
##          Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9970   0.9974   0.9951   0.9834   0.9926
## Specificity      0.9998   0.9981   0.9961   0.9984   1.0000
## Pos Pred Value    0.9994   0.9921   0.9817   0.9916   1.0000
## Neg Pred Value    0.9988   0.9994   0.9990   0.9968   0.9983
## Prevalence        0.2845   0.1935   0.1743   0.1638   0.1839
## Detection Rate    0.2836   0.1930   0.1735   0.1611   0.1825
## Detection Prevalence 0.2838   0.1946   0.1767   0.1624   0.1825
## Balanced Accuracy  0.9984   0.9977   0.9956   0.9909   0.9963

```

From the result above, it is clear that only few entries were predicted wrong (52 to be precise). The accuracy is also reported to be 99%, leaving the out-of-sample error at only 1%.

Conclusion and model implementation

The previous section clearly shows that the random forest algorithm generated a model which accuracy was the highest (99%). It is logical to implement this model to predict the classe on the testing data set.

```
pred_test <- predict(model_rf, data_test)
```

```
pred_test
```

```

## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E

```

```
table(pred_test)
```

```

## pred_test
## A B C D E
## 7 8 1 1 3

```

References

The data used in this project was obtained from this source: <http://groupware.les.inf.puc-rio.br/har>